

Tutorial - Spatial Analysis in R with Open Geodata

Egge-Jan Pollé - Tensing GIS Consultancy B.V.¹

Willy Tadema - Provincie Groningen²

Version 0.6.9 - August 22, 2018

Introduction

This is the manual for a tutorial session to be delivered on Wednesday 12 September 2018 at the 6th International Conference on the Use of R in Official Statistics (uRos2018) at the Dutch Office for National Statistics (CBS) in The Hague.

- Date: **Wednesday, September 12, 2018**
- Time: **9:00 - 12:30**
- Location: **Tinbergenzaal** (at CBS, Henri Faasdreef 312, 2492JP The Hague, The Netherlands)

Course material available on GitHub


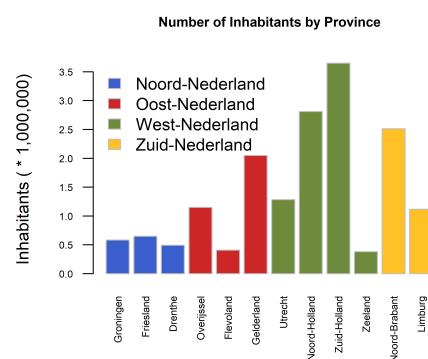
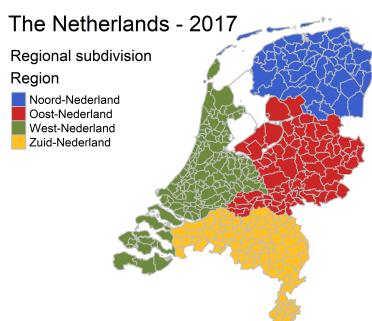
The URL of this repository: https://github.com/TWIAV/Spatial_Analysis_in_R_with_Open_Geodata

Prerequisites

- Attendees should bring their own laptop with (64-bit) **R**, and preferably also **RStudio**, installed.
- Shortly before the training session we will publish a list of additional packages (e.g. **sf**, **mapview** and **tmap**) to be installed - please keep an eye on our [GitHub repository](#).
- No specific prior experience (with R) is required for this introductory session, though some basic experience in one or more of the following fields is certainly helpful: Data Science, Programming/Software Development and/or GIS.

¹<https://www.linkedin.com/in/ejhpolle/>

²<https://www.linkedin.com/in/willytadema/>



Contents

Introduction	1
Prerequisites	1
Course material available on GitHub	1
Managing Geospatial Vector Data in R	4
1 The package sf (Simple Features for R)	4
1.1 st_as_sf(): converting a data.frame into an sf object	5
1.2 Interactive Viewing of Spatial Data in R	8
1.2.1 The package mapview	8
1.2.2 The package tmap	9
1.3 st_write()	10
1.3.1 Export data to an ESRI Shapefile	10
1.3.2 Export data to a GeoJSON file	11
Accessing Geospatial Vector Data over the Internet	12
2 Downloadable shapefiles	12
2.1 Download and unzip the shapefile	12
2.2 Load the shapefile	13
3 OGC Web Feature Service (WFS)	14
3.1 Introduction	14
3.2 WFS Outputformat: GML vs. GeoJSON	14
3.3 Access a WFS service: request=GetCapabilities	14
3.4 Retrieve data from a WFS service: request=GetFeature	15
3.4.1 request=GetFeature - an example from the Netherlands	15
3.4.2 request=GetFeature - an example from Finland	17
3.5 Get a description of a dataset: request=DescribeFeatureType	19
3.6 Additional parameters to the GetFeature request	19
3.6.1 Limit the number of records returned: the count parameter	19
3.6.2 Limit the number of columns returned: the PropertyName parameter	20
3.7 Further reading	20
4 ArcGIS REST Service	22
4.1 Introduction	22
4.2 Exploring data in the ArcGIS Living Atlas of the World	22
4.3 Query the Feature Service	22
4.4 Additional parameters to the query request	26
4.4.1 Filter the records returned: specify the where-clause	26
4.4.2 Limit the number of columns returned: the outFields parameter	26
4.5 Advanced topics	27
5 Statistics Netherlands (CBS) StatLine databank as open data	31
5.1 the package cbsodataR	31
5.2 Exercises with regional statistical data	32
5.2.1 Preparing the data	32
5.2.2 Extracting NUTS 1 and NUTS 2 regions	34
5.2.3 barplot()	36
5.2.4 pie()	38
5.2.5 Extracting municipal data	40
5.2.6 Merging sf and data.frame objects	41
5.2.7 qtm()	42
5.3 Factor: group municipalities into categories	43

6 Spatial Reference Systems	45
6.1 Coordinate systems: Geographic vs. Projected	45
6.2 epsg codes and. proj4string definitions	45
6.3 Reprojecting data	46
References	47
Appendix A: List of abbreviations used	48
Appendix B: Additional exercises	49
OGC Web Feature Service (WFS)	49
request=GetFeature - another example from the Netherlands	49

Managing Geospatial Vector Data in R

1 The package `sf` (Simple Features for R)

To manage spatial data in R in this manual we will be using the library `sf`, a relativle new addition to the R universe. The package was [released on CRAN in January 2017](#).

This package provides support for simple features, which is a standardized way to encode spatial vector data.

`sf` links directly to three important geospatial libraries, to unlock their power for use in R:

- GDAL for reading and writing data
- GEOS for geometrical operations
- Proj.4 for projection conversions and datum transformations.

The package `sf` on CRAN:

<https://cran.r-project.org/package=sf>

If the package is not yet installed, you can install it with the following command:

```
install.packages("sf")
```

To be able to manage your spatial data you will first have to load the package `sf`:

```
library(sf)
```

```
## Linking to GEOS 3.6.1, GDAL 2.2.3, proj.4 4.9.3
```

The real geospatial powers behind `sf`

- GDAL: the **Geospatial Data Abstraction Library** is a translator library for raster and vector geospatial data formats. Website: <http://www.gdal.org/>
- GEOS: the **Geometry Engine, Open Source** contains the complete functionality of the OpenGIS Simple Features for SQL spatial predicate functions and spatial operators. Website: <https://trac.osgeo.org/geos>
- Proj.4: **PROJ** is a generic coordinate transformation software, that transforms coordinates from one coordinate reference system (CRS) to another. This includes cartographic projections as well as geodetic transformations. Website: <http://proj4.org/>

1.1 st_as_sf(): converting a data.frame into an sf object

In this first exercise we will try to get some Dutch airports on the map. Let's start with defining some variables:

```
ap <- "Lelystad Airport"
cd <- "LEY"
class(cd)
```

```
## [1] "character"
lat <- 52.460278
lon <- 5.527222
class(lon)
```

```
## [1] "numeric"
```

Now we will combine these variables into a data.frame:

```
airport <- data.frame(ap, cd, lat, lon, stringsAsFactors = FALSE)
```

Wait, we have found a CSV file with some more airports:

```
NL_Airports <-
  read.csv("http://www.twiav.nl/files/NL_Airports.csv", stringsAsFactors = FALSE)
```

```
NL_Airports
```

```
##                                     airport iata latitude longitude
## 1  Amsterdam Airport Schiphol   AMS  52.30833  4.760833
## 2 Rotterdam The Hague Airport  RTM  51.95000  4.433333
## 3    Groningen Airport Eelde   GRQ  53.12500  6.583333
## 4      Eindhoven Airport     EIN  51.45028  5.374444
## 5 Maastricht Aachen Airport   MST  50.91583  5.776944
```

Let's add our single airport to this dataset:

```
NL_Airports <- rbind(NL_Airports, airport)
```

```
## Error in match.names(clabs, names(xi)): names do not match previous names
```

```
# Oh yeah, of course: this will generate an error...
```

We can easily solve this issue setting the names equal to eachother:

```
names(airport) <- names(NL_Airports)
```

Now we can rbind the two data.frames:

```
NL_Airports <- rbind(NL_Airports, airport)
NL_Airports
```

```
##                                     airport iata latitude longitude
## 1  Amsterdam Airport Schiphol   AMS  52.30833  4.760833
## 2 Rotterdam The Hague Airport  RTM  51.95000  4.433333
## 3    Groningen Airport Eelde   GRQ  53.12500  6.583333
## 4      Eindhoven Airport     EIN  51.45028  5.374444
## 5 Maastricht Aachen Airport   MST  50.91583  5.776944
## 6      Lelystad Airport     LEY  52.46028  5.527222
```

```
class(NL_Airports)
```

```
## [1] "data.frame"
```

Now, how do we convert this `data.frame` into an `sf` object? We have got information on the latitude and the longitude of our airports, and we can use this to create geometry, in this case: points.

Please note: the Coordinate Reference System (CRS) of these lat\lon coordinates is WGS84, which has been given the EPSG code 4326. In chapter 6 we will have a closer look at spatial reference systems.

Now we can use the function `st_as_sf()` to convert our `data.frame` - with longitude first:

```
library(sf)
```

```
NL_Airports <- st_as_sf(NL_Airports, coords = c("longitude","latitude"), crs = 4326)
```

```
class(NL_Airports)
```

```
## [1] "sf"           "data.frame"
```

Our dataset is now both a `data.frame` and an `sf` object, with all the conveniences which come with this dual status!

Figure 1: Our first `sf` object in the **RStudio Environment** pane

The screenshot shows the RStudio Environment pane. At the top, there are tabs for Environment, History, and Connections. Below the tabs, there are buttons for Import Dataset and Global Environment. The main area displays the contents of the `NL_Airports` object. It is a list containing two items: `airport` and `NL_Airports`. The `airport` item is a `data.frame` with 1 observation and 4 variables: `airport` (value: "Lelystad Airport"), `iata` (value: "LEY"), `latitude` (value: 52.5), and `longitude` (value: 5.53). The `NL_Airports` item is an `sf` object with 6 observations and 3 variables: `airport` (values: "Amsterdam Airport Schiphol", "Rotterdam The Hague Airport", "Groningen Airport E..."), `iata` (values: "AMS", "RTM", "GRQ", "EIN", ...), and `geometry` (value: sfc_POINT of length 6; first list element: 'XY' num 4.76 52.31). Below the object listing, there is a section titled "Values" showing the values for the variables: `ap` ("Lelystad Airport"), `cd` ("LEY"), `lat` (52.460278), and `lon` (5.527222).

And when we plot the data to the screen we can see the sheer beauty of our **Simple Features**: the original columns containing the lat\lon information have been used to create a column containing points. And this geometry column is stored next to the attribute data, in the very same `data.frame`. Wonderful, isn't it?

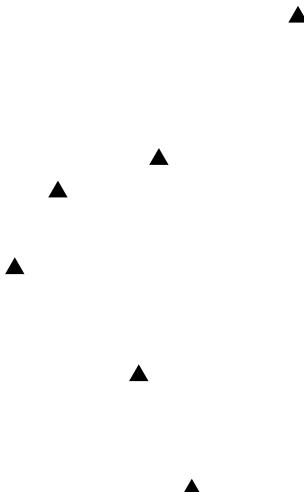
```
NL_Airports
```

```
## Simple feature collection with 6 features and 2 fields
## geometry type:  POINT
## dimension:      XY
## bbox:            xmin: 4.433333 ymin: 50.91583 xmax: 6.583333 ymax: 53.125
## epsg (SRID):    4326
## proj4string:    +proj=longlat +datum=WGS84 +no_defs
##                   airport iata          geometry
## 1   Amsterdam Airport Schiphol  AMS  POINT (4.760833 52.30833)
## 2   Rotterdam The Hague Airport  RTM  POINT (4.433333 51.95)
## 3   Groningen Airport Eelde     GRQ  POINT (6.583333 53.125)
## 4   Eindhoven Airport           EIN  POINT (5.374444 51.45028)
## 5   Maastricht Aachen Airport  MST  POINT (5.776944 50.91583)
## 6   Lelystad Airport           LEY  POINT (5.527222 52.46028)
```

And we can plot our airports like this:

```
plot(st_geometry(NL_Airports), main = "Airports in the Netherlands", pch = 17)
```

Airports in the Netherlands



1.2 Interactive Viewing of Spatial Data in R

In the previous paragraph we have seen a static map plotted on the **Plots** pane. But as a real Data Scientist we also want to be able to explore our data on an interactive map. Until a few years ago this would have meant switching back and forth between a desktop GIS and R. But in recent years some new packages have been developed to enable interactive map viewing in R.

In this paragraph we will look at two of these.

1.2.1 The package `mapview`

The package `mapview` on CRAN:

<https://cran.r-project.org/package=mapview>

If the package is not yet installed, you can install it with the following command:

```
install.packages("mapview")
```

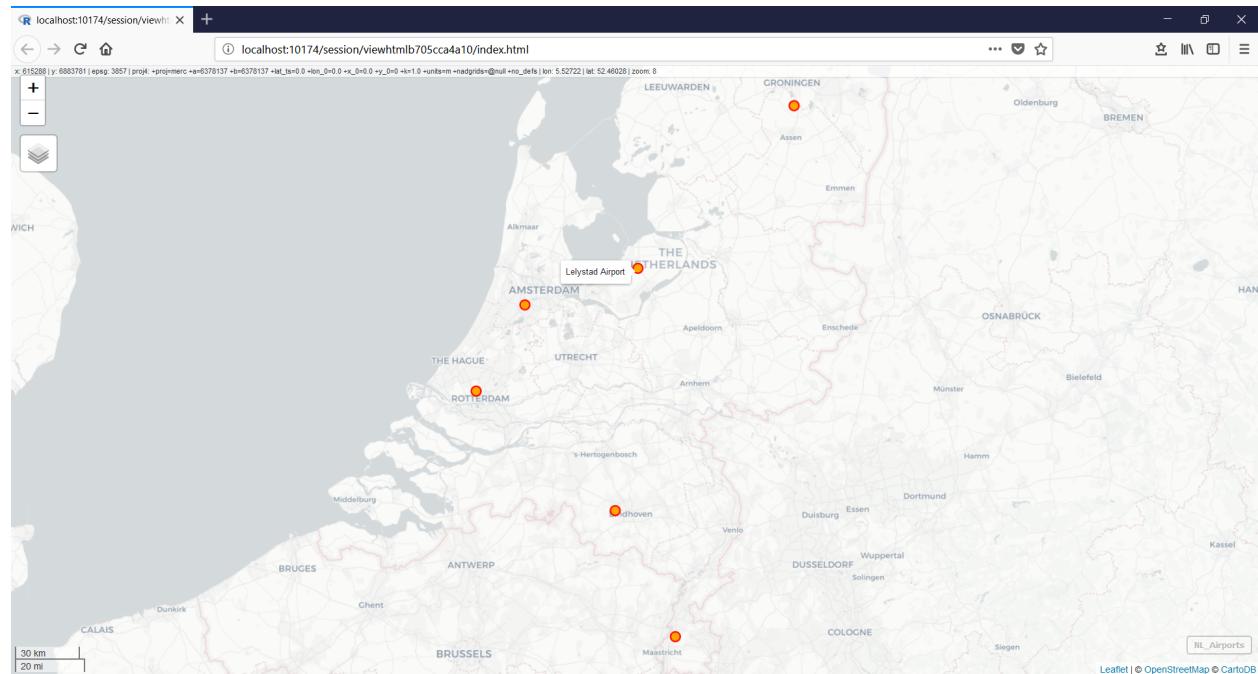
To be able to view your spatial data interactively you will first have to load the package `mapview`:

```
library(mapview)
```

Now we can open up an interactive map using a the function `mapview()` (see Figure 2):

```
mapview(NL_Airports, color = "red", col.regions = "orange", alpha.regions = 1, label =
  NL_Airports$airport)
```

Figure 2: An interactive map can be opened up using a simple `mapview()` statement



1.2.2 The package `tmap`

The package `tmap` on CRAN:

<https://cran.r-project.org/package=tmap>

If the package is not yet installed, you can install it with the following command:

```
install.packages("tmap")
```

To be able to view your spatial data interactively you will first have to load the package `tmap`:

```
library(tmap)
```

1.3 st_write()

To share the spatial data in a simple features object with non-R users we have to write it to a spatial database or file format with the function `st_write()`

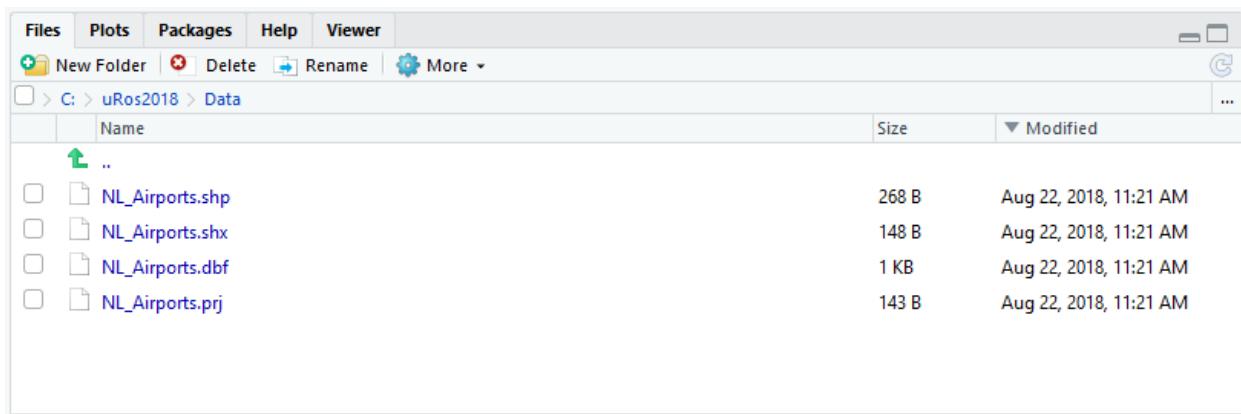
1.3.1 Export data to an ESRI Shapefile

A well-known format to write to is the ESRI Shapefile format (see Figure 3):

```
st_write(NL_Airports, "NL_Airports.shp")
```

```
## Writing layer `NL_Airports' to data source `NL_Airports.shp' using driver `ESRI Shapefile'
## features:       6
## fields:        2
## geometry type: Point
```

Figure 3: The shapefile with the corresponding files are written to your Working Directory



The ESRI Shapefile

The shapefile format is a well-known and still rather popular geospatial vector data format for Geographic Information System (GIS) software. It spatially describes vector features - points, lines, and polygons - with attribute data attached. The shapefile is a 'classic' GIS file format in the sense that it stores geometry and attribute data in separate files (in this chapter we will discover that a single shapefile in reality consists of multiple files) as opposed to more modern spatial database or file formats, where geometry and attributes are stored together in a single table or file.

The shapefile format has been developed by [Esri](#) and over the years has become a *de facto* standard for data interoperability among Esri and other GIS software products.

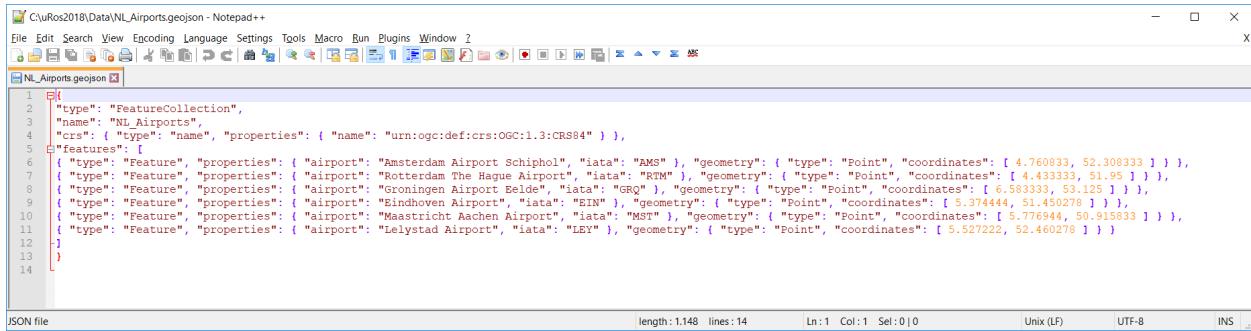
1.3.2 Export data to a GeoJSON file

Another format we may use is GeoJSON (see Figure 4):

```
st_write(NL_Airports, "NL_Airports.geojson")
```

```
## Writing layer `NL_Airports' to data source `NL_Airports.geojson' using driver `GeoJSON'
## features:       6
## fields:        2
## geometry type: Point
```

Figure 4: The GeoJSON file can be viewed in a text editor



The screenshot shows a Notepad++ window titled 'ChuRos2018\Data\NL_Airports.geojson - Notepad++'. The file contains a JSON object representing a FeatureCollection. The code is as follows:

```

1  {
2   "type": "FeatureCollection",
3   "name": "NL Airports",
4   "crs": { "type": "name", "properties": { "name": "urn:ogc:def:crs:OGC:1.3:CRS84" } },
5   "features": [
6     {
7       "type": "Feature",
8       "properties": { "airport": "Amsterdam Airport Schiphol", "iata": "AMS" },
9       "geometry": { "type": "Point", "coordinates": [ 4.760833, 52.308333 ] }
10      },
11      {
12       "type": "Feature",
13       "properties": { "airport": "Rotterdam The Hague Airport", "iata": "RTM" },
14       "geometry": { "type": "Point", "coordinates": [ 4.433333, 51.95 ] }
15      },
16      {
17       "type": "Feature",
18       "properties": { "airport": "Groningen Airport Eelde", "iata": "GRQ" },
19       "geometry": { "type": "Point", "coordinates": [ 6.583333, 53.125 ] }
20      },
21      {
22       "type": "Feature",
23       "properties": { "airport": "Eindhoven Airport", "iata": "EIN" },
24       "geometry": { "type": "Point", "coordinates": [ 5.374444, 51.450278 ] }
25      },
26      {
27       "type": "Feature",
28       "properties": { "airport": "Maastricht Aachen Airport", "iata": "MST" },
29       "geometry": { "type": "Point", "coordinates": [ 5.776944, 50.915833 ] }
30      },
31      {
32       "type": "Feature",
33       "properties": { "airport": "Lelystad Airport", "iata": "LEY" },
34       "geometry": { "type": "Point", "coordinates": [ 5.527222, 52.460278 ] }
35     }
36   ]
37 }
```

The status bar at the bottom indicates: length:1.148 lines:14 Ln:1 Col:1 Sel:0|0 Unix (LF) UTF-8 INS .

Accessing Geospatial Vector Data over the Internet

2 Downloadable shapefiles

In this chapter we will learn how to download, to unzip and to load shapefiles **using R**. We have seen the ESRI Shapefile before in paragraph 1.3.1. There will be no need to use your web browser, your file explorer or a zip utility - the whole process can be completed using just a few lines of R code.

It is not uncommon for public organizations - including national statistical institutes - to distribute geographic information in this shapefile format. So, you will find shapefiles on the [page with geographic data of the Dutch Office for National Statistics \(CBS\)](#) and also on the [STATISTIK AUSTRIA open.data Portal](#).



2.1 Download and unzip the shapefile

```
# Store the URL to the file to download in a variable
URL2zip <- "http://data.statistik.gv.at/data/OGDEXT_GEM_1_STATISTIK_AUSTRIA_20180101.zip"

# Create a temporary file
zip_file <- tempfile(fileext = ".zip")

# Download the file
download.file(URL2zip, destfile = zip_file, mode = "wb")

# Create a subfolder in your working directory to store the unzipped data
dir.create("./Data", showWarnings = FALSE)

# Unzip the file
unzip(zip_file, exdir = "./Data")

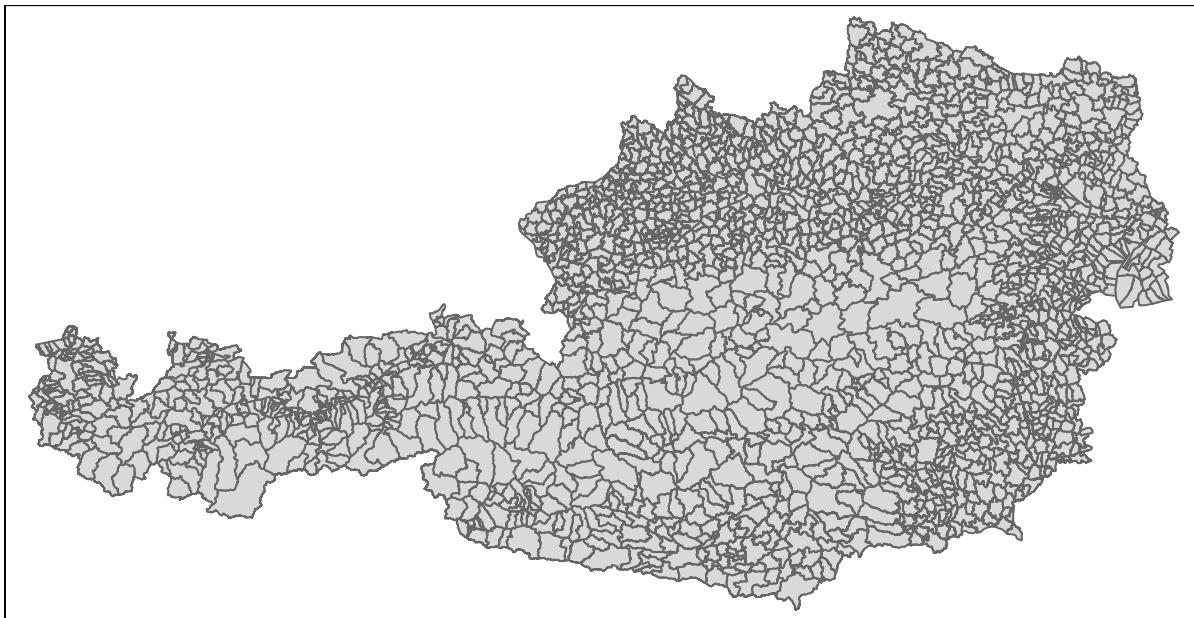
# After unzipping you can delete (i.e. unlink) the file
unlink(zip_file)

# Remove variables you do not longer need
rm(URL2zip, zip_file)
```

2.2 Load the shapefile

```
library(sf)
library(tmap)
AUSTRIA_GEM_20180101 <- st_read("./Data/STATISTIK_AUSTRIA_GEM_20180101.shp")

## Reading layer `STATISTIK_AUSTRIA_GEM_20180101' from data source `C:\temp\git\Spatial_Analysis_in_R_w...
## Simple feature collection with 2120 features and 2 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:            xmin: 112518.2 ymin: 275472 xmax: 685444.5 ymax: 570431.1
## epsg (SRID):    31287
## proj4string:    +proj=lcc +lat_1=49 +lat_2=46 +lat_0=47.5 +lon_0=13.33333333333333 +x_0=400000 +y_0=1.3145,-5.393,-2.3887 +units=m +no_defs
qtm(AUSTRIA_GEM_20180101)
```



3 OGC Web Feature Service (WFS)

3.1 Introduction

In this chapter we will learn how **to use R as a client to access data using a Web Feature Service (WFS)**. WFS is a Data Access Standard which is defined and maintained by the Open Geospatial Consortium (OGC). The WFS Interface Standard defines a set of interfaces for accessing geographic information over the Internet. It offers the means to retrieve geographic features and their properties through a highly configurable interface and in a manner independent of the underlying data stores they publish.

The standard is used - both in the public and private sector and

in the academic world - to publish vector geospatial datasets in a way that makes it easy for receiving organisations to conduct analysis on the data supplied.

In a short chapter like this it will not be possible to cover all ins and outs of the WFS Interface Standard. The main goal here is to get you up and running and to whet your appetite for more. It might all look a bit technical and intimidating at first, but as soon as you have the syntax of your request right, you will be able to retrieve valuable spatial data in the blink of an eye.

3.2 WFS Outputformat: GML vs. GeoJSON

By default, a WFS returns data in Geography Markup Language (GML) which is written as eXtensible Markup Language (XML). However, many WFS services also offer the option to request the output in GeoJSON, a geospatial data interchange format based on JavaScript Object Notation (JSON).

By its very nature, GML data is difficult to process, because - as with most XML based grammars - there are two parts to the grammar: the schema that describes the document and the instance document that contains the actual data.

On the contrary, the GeoJSON standard clearly defines several types of JSON objects and the manner in which they are combined to represent data about geographic features, their properties, and their spatial extents.

In general **GDAL**, the translator library behind the `sf` functions `st_read()` and `st_write` (see chapter 1), gives better results with GeoJSON as opposed to GML.

So, in the exercises in this manual, when accessing a WFS service to retrieve data, we will always add the parameter `outputFormat=application/json`.

3.3 Access a WFS service: `request=GetCapabilities`

In general, organisations publishing data using WFS, will provide you with a URL to their WFS server which also contains a `GetCapabilities` request. Once you know the **capabilities** of the service (i.e. once you know what is on offer) you can start building your own request to the server.

In the exercises below we will use data on the municipal division in the Netherlands. These data are offered by the host of this conference, the Dutch statistical office, through [Publieke Dienstverlening op de Kaart](#) (PDOK) and the [Nationaal Georegister](#) (NGR).

Some background information on the dataset used, can be found here: [Dataset: CBS Gebiedsindelingen](#).

The actual URL with the `GetCapabilities` request is:

<https://geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs?request=GetCapabilities>

When you click this link the server response will be shown in your browser in XML format.

We are not going to study this XML response line-by-line. Not now. But please leave the web page open in your browser for later reference.

3.4 Retrieve data from a WFS service: request=GetFeature

3.4.1 request=GetFeature - an example from the Netherlands

A WFS server responding to a *GetFeature* request returns a collection of geographic feature instances filtered according to a criteria set by the requesting client. We will start with a simple request to download a full feature collection without any constraints to filter the content by. The GetFeature request queries the server with a set of parameters, which are concatenated to the URL with an *ampersand* (&).

In the script below we use the `httr` package. This allows us to store the different parameters of our request in a list, only to build the full URL at the end. We do so for readability reasons and to allow for easy modification of our request at a later stage if necessary. And the function `build_url()` will return a properly encoded URL.

A WFS service can offer one or more feature collections, see the `<FeatureTypeList>` section in the XML response to the `GetCapabilities` request. The service we are accessing here offers quite some feature collections, i.e. multiple regional divisions for the years 1995 up to the current year. What we are interested in here, is the municipal division for the year 2017. After some browsing through the XML response, we have found a `<FeatureType>` with the `<Name> cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd`. And that's the value we are giving to the `typename` parameter. (This `typename` parameter determines the collection of feature instances to return.)

Also, do not forget to add a parameter to ask for output in GeoJSON format (as discussed in paragraph 3.2).

This is the script to populate the full request:

```
# NL: Example with Dutch data
library(sf)
library(tmap)
library(httr)
library(data.table)

url <- list(hostname = "geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs",
            scheme = "https",
            query = list(service = "WFS",
                         version = "2.0.0",
                         request = "GetFeature",
                         typename =
                           "cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd",
                         outputFormat = "application/json")) %>%
  setattr("class","url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

Now we want to feed this response directly into R, like this:

```
NL_Municipalities2017 <- st_read(request)
```

```
## Reading layer `OGRGeoJSON` from data source `https://geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs?request=GetFeature&typename=cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd&outputFormat=application/json`
## Simple feature collection with 388 features and 5 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:            xmin: 13565.4 ymin: 306846.9 xmax: 277992.8 ymax: 619291
## epsg (SRID):    28992
## proj4string:    +proj=sterea +lat_0=52.15616055555555 +lon_0=5.38763888888889 +k=0.9999079 +x_0=1550000 +y_0=406857 +towgs84=1550000,406857,-1.87033 +units=m +no_defs
```

Have a look at the result:

```
qtm(NL_Municipalities2017)
```



```
head(NL_Municipalities2017)
```

```
## Simple feature collection with 6 features and 5 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:            xmin: 226868.9 ymin: 562754.2 xmax: 276821.5 ymax: 604994.7
## epsg (SRID):    28992
## proj4string:    +proj=sterea +lat_0=52.1561605555555 +lon_0=5.38763888888889 +k=0.9999079 +x_0=15500
##               0.406857,0.350733,-1.87035,4.0812 +units=m +no_defs
##                                     id statcode
## 1 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_765c GM0003
## 2 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_765d GM0005
## 3 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_765e GM0007
## 4 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_765f GM0009
## 5 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_7660 GM0010
## 6 cbs_gemeente_2017_gegeneraliseerd.fid-3aa9a616_165621cca67_7661 GM0014
##           statnaam jrstatcode rubriek      geometry
## 1     Appingedam 2017GM0003 gemeente MULTIPOLYGON (((254580.7 59...
## 2          Bedum 2017GM0005 gemeente MULTIPOLYGON (((235432.1 59...
## 3   Bellingwedde 2017GM0007 gemeente MULTIPOLYGON (((276518.1 56...
## 4        Ten Boer 2017GM0009 gemeente MULTIPOLYGON (((245194.7 59...
## 5       Delfzijl 2017GM0010 gemeente MULTIPOLYGON (((262016.6 58...
## 6      Groningen 2017GM0014 gemeente MULTIPOLYGON (((243244.4 58...
```

3.4.2 request=GetFeature - an example from Finland

In this paragraph we will retrieve data from [Tilastokeskus \(Statistics Finland\)](#). In the English section of their website we found a nice dataset: [Population by municipality-based units](#).

And from this map it is especially the layer **Population 2017 by municipalities 2018 (kunta_vaki2017)** we would like to investigate.

We will follow the same steps as in the Dutch example above. For an explanation of the steps taken, please refer to paragraph 3.3 and 3.4.1.



The URL with GetCapabilities request is:

<http://geo.stat.fi/geoserver/vaestoa/value/wfs?request=GetCapabilities>

In this XML response - in the <FeatureTypeList> section - we have found a <FeatureType> with the <Name> **vaestoa/value:kunta_vaki2017**. That's the value we are giving to the typename parameter.

The script to build the full URL is similar to the one before - the only parameters we have modified are the hostname and the typename:

```
# FI: Example with Finnish data
library(sf)
library(tmap)
library(httr)
library(data.table)

url <- list(hostname = "geo.stat.fi/geoserver/vaestoa/value/wfs",
            scheme = "https",
            query = list(service = "WFS",
                         version = "2.0.0",
                         request = "GetFeature",
                         typename = "vaestoa/value:kunta_vaki2017",
                         outputFormat = "application/json")) %>%
  setattr("class", "url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

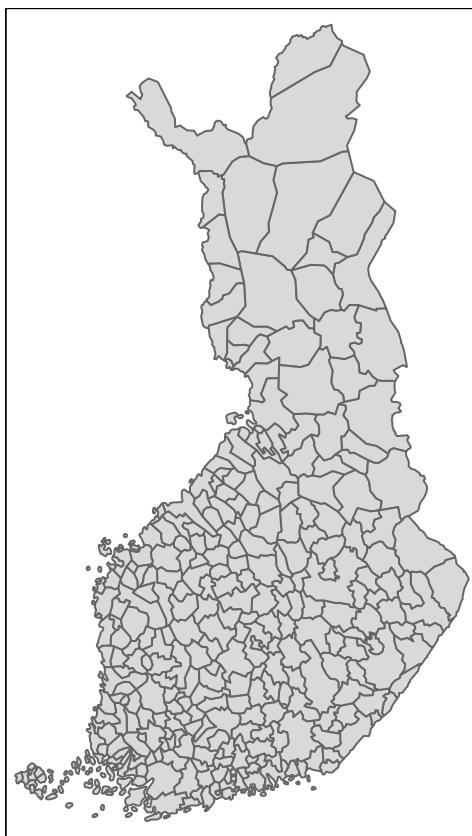
Now we want to feed this response directly into R, like this:

```
FI_Municipalities2018_Pop2017 <- st_read(request)
```

```
## Reading layer `OGRGeoJSON` from data source `https://geo.stat.fi/geoserver/vaestoa/value/wfs/?service=WFS&version=2.0.0&request=GetFeature&typename=vaestoa/value:kunta_vaki2017&outputFormat=application/json`
## Simple feature collection with 311 features and 20 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:            xmin: 83747.59 ymin: 6637032 xmax: 732907.8 ymax: 7776431
## epsg (SRID):   3067
## proj4string:    +proj=utm +zone=35 +ellps=GRS80 +towgs84=0,0,0,0,0,0,0 +units=m +no_defs
```

Have a look at the result:

```
qtm(FI_Municipalities2018_Pop2017)
```



Optionally, you can have a look at the attribute data in the **Data Viewer**: `View(FI_Municipalities2018_Pop2017)`

With some basic knowledge of the Finnish language you should now be able to calculate that Finland had more or less 5.5 million inhabitants in 2017 - of which more or less 50 percent are males, and the other half females:

Finnish for Data Scientists: some useful words

- **Miehet** = Men
- **Naiset** = Women
- **Nimi** = Name
- **Vaesto** = Population

```
sum(FI_Municipalities2018_Pop2017$vaesto)
```

```
## [1] 5513130
```

```
sum(FI_Municipalities2018_Pop2017$miehet)
```

```
## [1] 2719131
```

```
sum(FI_Municipalities2018_Pop2017$naiset)
```

```
## [1] 2793999
```

3.5 Get a description of a dataset: `request=DescribeFeatureType`

In the previous exercises we immediately executed GetFeature requests to WFS services - mainly because the trainer told us to do so - without actually knowing anything about these datasets.

To get a description of the dataset, you can execute a DescribeFeatureType request first. This returns a description - in XML format - of the structure, including properties, of the feature type specified in the request.

For the Dutch municipalities dataset this request would look like this: https://geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs?service=WFS&version=2.0.0&request=DescribeFeatureType&typename=cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd

And for the Finnish dataset like this: http://geo.stat.fi/geoserver/vaestotalue/wfs?service=WFS&version=2.0.0&request=DescribeFeatureType&typename=vaestotalue:kunta_vaki2017

3.6 Additional parameters to the GetFeature request

Additional parameters can be added to a GetFeature request to further filter or convert the response from the WFS.

To include additional parameters to a request, simply add an *ampersand* (&) at the end of the URL, then add the name of the parameter, an equal sign (=) and the value to assign to the parameter. Of course we will not do this manually; we will use the `httr` package to build this URL with additional parameters for us.

Below we will discuss a few of the parameters available.

3.6.1 Limit the number of records returned: the `count` parameter

To have a look at the structure of a dataset - before retrieving the full dataset - you can choose to limit the number of records returned with the `count` parameter.

The GetFeature request below limits the number of Finnish municipalities returned to only 5:

```
# FI: Example with Finnish data
library(httr)
library(data.table)

url <- list(hostname = "geo.stat.fi/geoserver/vaestotalue/wfs",
            scheme = "https",
            query = list(service = "WFS",
                         version = "2.0.0",
                         request = "GetFeature",
                         typename = "vaestotalue:kunta_vaki2017",
                         count = 5,
                         outputFormat = "application/json")) %>%
  setattr("class", "url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

3.6.2 Limit the number of columns returned: the `PropertyName` parameter

To restrict a GetFeature request by attribute rather than feature, use the `PropertyName` parameter. You can specify a single attribute, or multiple attributes separated by commas.

The GetFeature request below only returns the geometry and the columns *nimi* (Name) and *vaesto* (total population) for the Finnish municipalities:

```
# FI: Example with Finnish data
library(httr)
library(data.table)

url <- list(hostname = "geo.stat.fi/geoserver/vaestovalue/wfs",
            scheme = "https",
            query = list(service = "WFS",
                          version = "2.0.0",
                          request = "GetFeature",
                          typename = "vaestovalue:kunta_vaki2017",
                          propertynames = "geom,nimi,vaesto",
                          outputFormat = "application/json")) %>%
  setattr("class", "url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

3.7 Further reading

The OGC offers an official tutorial module - [OGC E-learning](#) - covering it's activities and the several standards it maintains. The specific chapter on [WFS](#) can be found [here](#).

Figure 5: On the overview page metadata about the Feature Layer USA States (Generalized) is presented

The screenshot shows the ArcGIS overview page for the 'USA States (Generalized)' feature layer. At the top, there's a navigation bar with links for 'ArcGIS', 'Prijsen', 'Kaart', 'Scene', and 'Help'. A search bar and a login button ('Aanmelden') are also present. The main title is 'USA States (Generalized)'. Below it, there are three tabs: 'Overzicht' (selected), 'Gegevens', and 'Visualisatie'. The 'Overzicht' tab contains a small map of the United States with orange states, a brief description, and a creation date of 'Gemaakt: 29 jun. 2013'. To the right, there are four buttons: 'Openen in Map Viewer', 'Openen in Scene Viewer', 'Openen in ArcGIS Desktop', and 'Metadata'. The 'Gegevens' tab is visible but not selected. The 'Visualisatie' tab is also visible but not selected.

Beschrijving

USA States (Generalized) provides 2017 boundaries for the States of the United States in the 50 states and the District of Columbia. The linework has been generalized for increased performance and best viewed at smaller scales.

Attribute fields include estimated 2017 total population, 2010 U.S. Census demographic information, and 2012 Census of Agriculture information for the USA States.

Kaartlagen

Details

Bron: Feature Service
Meest recente data-update: 18 mei 2018
18:52:31
Grootte: 656 KB

★★★★☆



Eigenaar

Figure 6: Here the Feature Layer USA States (Generalized) is shown in the ArcGIS Online Map Viewer

The screenshot shows the ArcGIS Online Map Viewer interface. The top navigation bar includes 'ArcGIS - USA States (Generalized)', a search bar, and buttons for 'Kaart wijzigen' and 'Aanmelden'. The main content area displays a world map with the USA highlighted in orange. The 'Content' pane on the left lists 'USA States (Generalized)' under 'Over' and 'Topografisch'. A legend is also visible. Below the map is a table titled 'USA States (Generalized) (Objecten: 51, Geselecteerd: 0)'. The table contains the following data:

STATE_NAME	Abbreviation	Region	POP2010	POP10_SQMI	WHITE	BLACK	AMERI_ES
Alaska	AK	Pacific	710.231	1,20	473.576	23.263	104.871
California	CA	Pacific	37.253.956	235,60	21.453.934	2.299.072	362.801
Hawaii	HI	Pacific	1.360.301	211,60	336.599	21.424	4.164
Md-hn	MD	Mountain	1.567.597	19,80	1.306.497	9.910	21.441

4 ArcGIS REST Service

4.1 Introduction

In this chapter we will learn how **to use R as a client to access data using an ArcGIS REST Service**. We will be accessing ArcGIS Server to retrieve Feature Layers using the [ArcGIS REST API](#). This API is part of ArcGIS Online, the “complete SaaS mapping platform” by the American GIS company [Esri](#). And as part of this platform Esri has launched the [ArcGIS Living Atlas of the World](#)³

Some datasets in this Living Atlas are Open Data. Some, but

by no means all, as many datasets are only accessible via an ArcGIS Online organizational account. But in cases where the source data is already Open Data, Esri publishes the data under the same conditions.



4.2 Exploring data in the ArcGIS Living Atlas of the World

After some browsing through the content of the Living Atlas we have found a nice (and open) dataset for this exercise: the **Feature Layer USA States (Generalized)**. An overview page with metadata for the dataset we have chosen can be found via this [link](#) (see Figure 5).

Before actually importing the data into R, you can explore the dataset - both the geometry and the attribute data - in the [ArcGIS Online Map Viewer](#) (see Figure 6).

4.3 Query the Feature Service

On the overview page we have found the Service URL pointing to the layer **USA_States_Generalized** in the Services Directory (see Figure 7):

https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0

One of the **Supported Operations** for this layer is **Query**. The **Query** page (see Figure 8) offers an interface to construct the query:

https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0/query

Let's try to build a request to retrieve the full dataset:

- The first parameter is **where**. It is obligatory and cannot be omitted. If there is no *where-clause* (i.e. if you want to extract the full dataset) you have to populate this parameter with the value **1=1** (which is - of course - always true...)
- To get all the attribute columns you have to give the parameter **outFields** the value *****
- To retrieve the spatial data the parameter **returnGeometry** should be set to **true**
- And the **outputformat** - **f** - is set to **geojson**

³Some background information on the Living Atlas of the World can be found in this [article](#), where Esri also invites YOU to “*Become a user and a contributor*”.

Figure 7: On this page the **Feature Layer USA_States_Generalized** is presented in the **ArcGIS REST Services Directory**

The screenshot shows a web browser window with the title bar "USA States (Generalized) > Layer: USA_States_Generalized". The address bar contains the URL "https://services.arcgis.com/P3ePLMys2RVChkjx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0". The main content area is titled "Layer: USA_States_Generalized (ID:0)". It displays various properties of the feature layer:

- Name:** USA_States_Generalized
- Type:** Feature Layer
- Geometry Type:** esriGeometryPolygon
- Description:** (empty)
- Copyright Text:** (empty)
- Min. Scale:** 0
- Max. Scale:** 0
- Default Visibility:** true
- Max Record Count:** 2000
- Supported query Formats:** JSON
- Use Standardized Queries:** True
- Extent:**
 - XMin: -19839092.3042881
 - YMin: 2145729.67991779
 - XMax: -7454985.14651674
 - YMax: 11542624.9160411
 - Spatial Reference: 102100

Figure 8: On this page you can enter parameters to query the **Feature Layer USA_States_Generalized**

The screenshot shows a web browser window with the title bar "USA States (Generalized) > Query: USA_States_Generalized". The address bar contains the URL "https://services.arcgis.com/P3ePLMys2RVChkjx/arcgis/rest/services/USA_States_Generalized/featureServer/0/query". The main content area is titled "Query: USA_States_Generalized (ID: 0)". It contains a form for querying the feature layer:

Where:	<input type="text"/>
Object IDs:	<input type="text"/>
Time:	<input type="text"/>
Input Geometry:	<input type="text"/>
Geometry Type:	Envelope
Input Spatial Reference:	<input type="text"/>
Spatial Relationship:	Intersects
Result Type:	None
Distance:	0.0
Units:	Meters
Return Geodetic:	<input checked="" type="radio"/> True <input type="radio"/> False
Out Fields:	<input type="text"/>
Return Geometry:	<input checked="" type="radio"/> True <input type="radio"/> False
Return Centroid:	<input checked="" type="radio"/> True <input type="radio"/> False
Geometry MultiPatch Option:	xyFootprint
Max Allowable Offset:	<input type="text"/>
Geometry Precision:	<input type="text"/>
Output Spatial Reference:	<input type="text"/>

This is the script to populate the full request:

```
# USA: An example with American data
library(sf)
library(tmap)
library(httr)
library(data.table)

url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = "USA_States_Generalized/FeatureServer/0/query",
            query = list(where = "1=1",
                         outFields = "*",
                         returnGeometry = "true",
                         f = "geojson")) %>%
  setattr("class", "url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

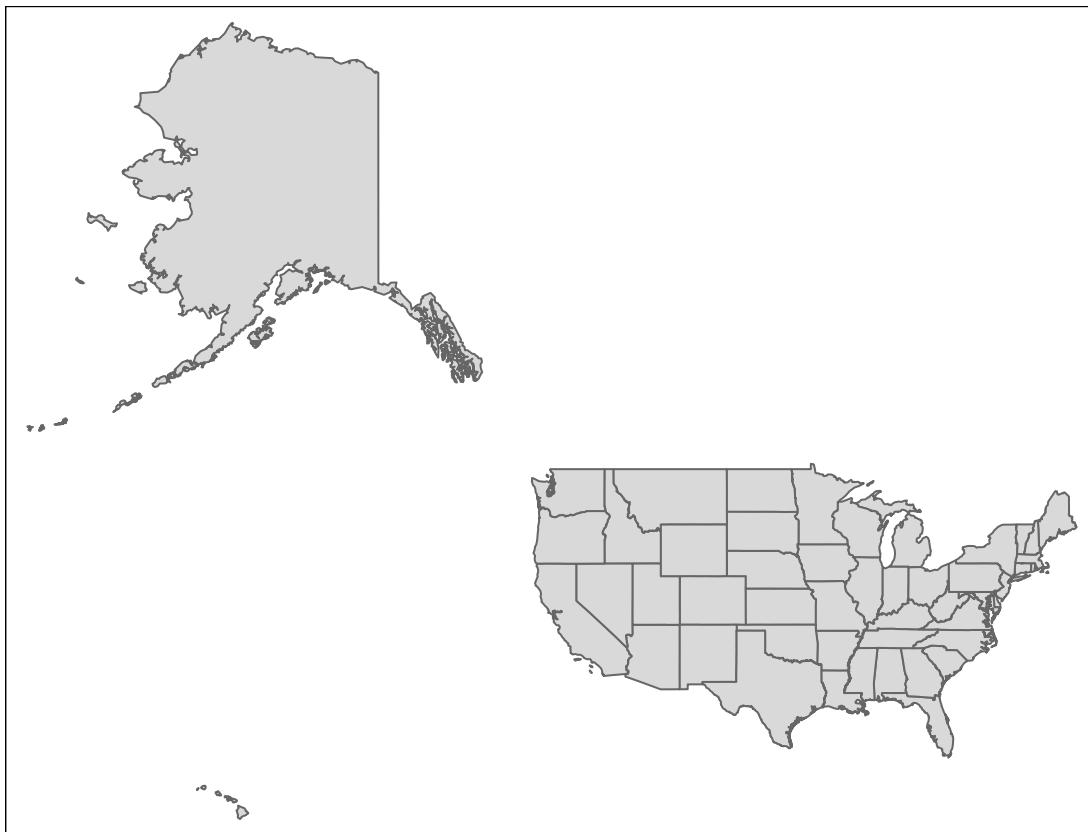
Now we want to feed this response directly into R, like this:

```
USA_States_2017 <- st_read(request)
```

```
## Reading layer `OGRGeoJSON' from data source `https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0'
## Simple feature collection with 51 features and 55 fields
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:           xmin: -19839090 ymin: 2145730 xmax: -7454985 ymax: 11542620
## epsg (SRID):   3857
## proj4string:   +proj=merc +a=6378137 +b=6378137 +lat_ts=0.0 +lon_0=0.0 +x_0=0.0 +y_0=0 +k=1.0 +units=m +no_defs
```

Have a look at the result:

```
qtm(USA_States_2017)
```



Anchorage

"Hey Chel you know it's kinda funny
Texas always seems so big
But you know you're in the largest State in the Union
When you're anchored down in Anchorage"

Michelle Shocked - Short Sharp Shocked - 1988

4.4 Additional parameters to the query request

4.4.1 Filter the records returned: specify the where-clause

To filter records you can specify a where-clause.

The query request below only returns those States which have more than 10 million inhabitants:

```
# USA: An example with American data
library(httr)
library(data.table)

url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = "USA_States_Generalized/FeatureServer/0/query",
            query = list(where = "POPULATION>10000000",
                         outFields = "STATE_NAME,POPULATION",
                         returnGeometry = "true",
                         f = "geojson")) %>%
  setattr("class","url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

4.4.2 Limit the number of columns returned: the outFields parameter

If you do not need all attribute columns of a dataset in your analysis, you can limit the number of columns returned by specifying the `outFields` parameter. You can specify a single attribute, or multiple attributes separated by commas.

The query request below only returns the geometry and the columns STATE_NAME and POPULATION for the United States:

```
# USA: An example with American data
library(httr)
library(data.table)

url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = "USA_States_Generalized/FeatureServer/0/query",
            query = list(where = "1=1",
                         outFields = "STATE_NAME,POPULATION",
                         returnGeometry = "true",
                         f = "geojson")) %>%
  setattr("class","url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

4.5 Advanced topics

So what exactly are the capabilities of a hosted feature service? To answer this question we can take a look at the information on the service in the ArcGIS REST Services Directory.

Let's give it a try. We'll query the [USA Rail Road](#) dataset in ESRI's Living Atlas of the World.

The service URL is https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_Railroads_1/FeatureServer.

A request to this URL will return the capabilities of service:

```
# USA: An example with USA Railroads data
```

```
library(httr)
library(data.table)
library(dplyr)
library(jsonlite)
```

```
url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = "USA_Railroads_1/FeatureServer",
            query = list(f = "json")) %>%
  setattr("class","url")
response <- build_url(url) %>% fromJSON()
```

Notice that the format of the response is json instead of geojson (query parameter f = "json").

A hosted feature service can have multiple layers. Every layer has a name and an id. The id is necessary for constructing the appropriate URL to fetch the features.

The response contains the layer information needed:

```
response$layers %>% select(id, name)
```

```
##   id      name
## 1  0 USA Railroads
```

The USA Rail Road hosted feature service only has one layer with called *USA Railroads*.

```
layer_id <- response$layers %>% filter(name == 'USA Railroads') %>% select(id)
layer_id
```

```
##   id
## 1  0
```

The id of the USA Railroads layer is 0. We use this information to construct the correct value for the path variable.

```
layer_path <- paste("USA_Railroads_1/FeatureServer", layer_id, sep = "/")
layer_path
```

```
## [1] "USA_Railroads_1/FeatureServer/0"
```

The response also other valuable like the maximum number of records returned per request:

```
response$maxRecordCount
```

```
## [1] 1000
```

We can get even more information on the capabilities of the railroad layer:

```
url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = layer_path,
            query = list(f = "json")) %>%
  setattr("class", "url")
response <- build_url(url) %>% fromJSON()
```

For example information on the fields in the dataset:

```
response$fields %>% select(name, type)

##          name           type
## 1    OBJECTID      esriFieldTypeOID
## 2     FRA_ID      esriFieldTypeInteger
## 3   FRA_REGION esriFieldTypeSmallInteger
## 4      STATE      esriFieldTypeString
## 5 SUBDIVISIO      esriFieldTypeString
## 6   RR_OWNER      esriFieldTypeString
## 7     TRACKS esriFieldTypeSmallInteger
## 8    NET_CODE      esriFieldTypeString
## 9    NET_DESC      esriFieldTypeString
## 10  PASSENGER      esriFieldTypeString
## 11     MILES      esriFieldTypeDouble
```

So now we know the query URL, the maximum number of records returned per request and the fieldnames in the dataset. We are ready to request the records!

But wait, all the railroads in the United States of America, wouldn't that be *a lot* of records? Let's find out, before we fetch them all.

We can do this by specifying the query parameter `returnCountOnly = "true"`:

```
url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = paste(layer_path, "query", sep = "/"),
            query = list(where = "1=1",
                         returnCountOnly = "true",
                         f = "geojson")) %>%
  setattr("class", "url")
request <- build_url(url)
response <- build_url(url) %>% fromJSON()
hits <- response$properties$count
hits
```

```
## [1] 173114
```

173114 is way more than the maximum record count per request of 1000!

Let's do a request to fetch the records.

```
url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = paste(layer_path, "query", sep = "/"),
            query = list(where = "1=1",
                         outFields = "OBJECTID, NET_DESC",
                         returnGeometry = "true",
                         f = "geojson")) %>%
  setattr("class", "url")
```

```

request <- build_url(url)

railroads <- st_read(request)

## Reading layer `OGRGeoJSON' from data source `https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0/query?where=1=1&geometryType=esriGeometryLine&inSR=4326&spatialRel=esriSpatialRelIntersects&outFields=OBJECTID,NET_DESC&returnGeometry=true&f=geojson'
## Simple feature collection with 173114 features and 2 fields
## geometry type:  LINESTRING
## dimension:      XY
## bbox:            xmin: -158.1522 ymin: 20.88018 xmax: -66.98391 ymax: 64.92615
## epsg (SRID):    4326
## proj4string:    +proj=longlat +datum=WGS84 +no_defs

```

What is the number of records in the railroads dataframe? One would expect, no more than 1,000.

```
nrow(railroads)
```

```
## [1] 173114
```

Isn't that strange? The dataframe contains all 173114 records. How is this possible?

Well the GDAL library makes it very easy for us. It detects that multiple requests are needed to fetch all the records and does this automagically for us. So no need to think of response paging, GDAL already has the solution for that!

A lot of the railroads are labelled 'abandonned'. Of course you can do this *after* you fetched the records, but this isn't very efficient. Let's change the parameter value of where in our request, so only exploited rail roads are included in the response.

```

url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = paste(layer_path, "query", sep = "/"),
            query = list(where = "NET_DESC <> 'Abandoned'",
                         outFields = "OBJECTID, NET_DESC",
                         returnGeometry = "true",
                         f = "geojson")) %>%
  setattr("class", "url")
request <- build_url(url)

exploited_railroads <- st_read(request)

## Reading layer `OGRGeoJSON' from data source `https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services/USA_States_Generalized/FeatureServer/0/query?where=1=1&geometryType=esriGeometryLine&inSR=4326&spatialRel=esriSpatialRelIntersects&outFields=OBJECTID,NET_DESC&returnGeometry=true&f=geojson'
## Simple feature collection with 134436 features and 2 fields
## geometry type:  LINESTRING
## dimension:      XY
## bbox:            xmin: -158.1522 ymin: 20.88018 xmax: -67.26739 ymax: 64.92615
## epsg (SRID):    4326
## proj4string:    +proj=longlat +datum=WGS84 +no_defs

```

```
nrow(exploited_railroads)
```

```
## [1] 134436
```

134436 is still a lot of records, but let's have a look at the result and add the railroads on top of a map of the USA:

```

url <- list(hostname = "services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/rest/services",
            scheme = "https",
            path = "USA_States_Generalized/FeatureServer/0/query",
            query = list(where = "1=1",
                         returnGeometry = "true",
                         f = "geojson")) %>%

```

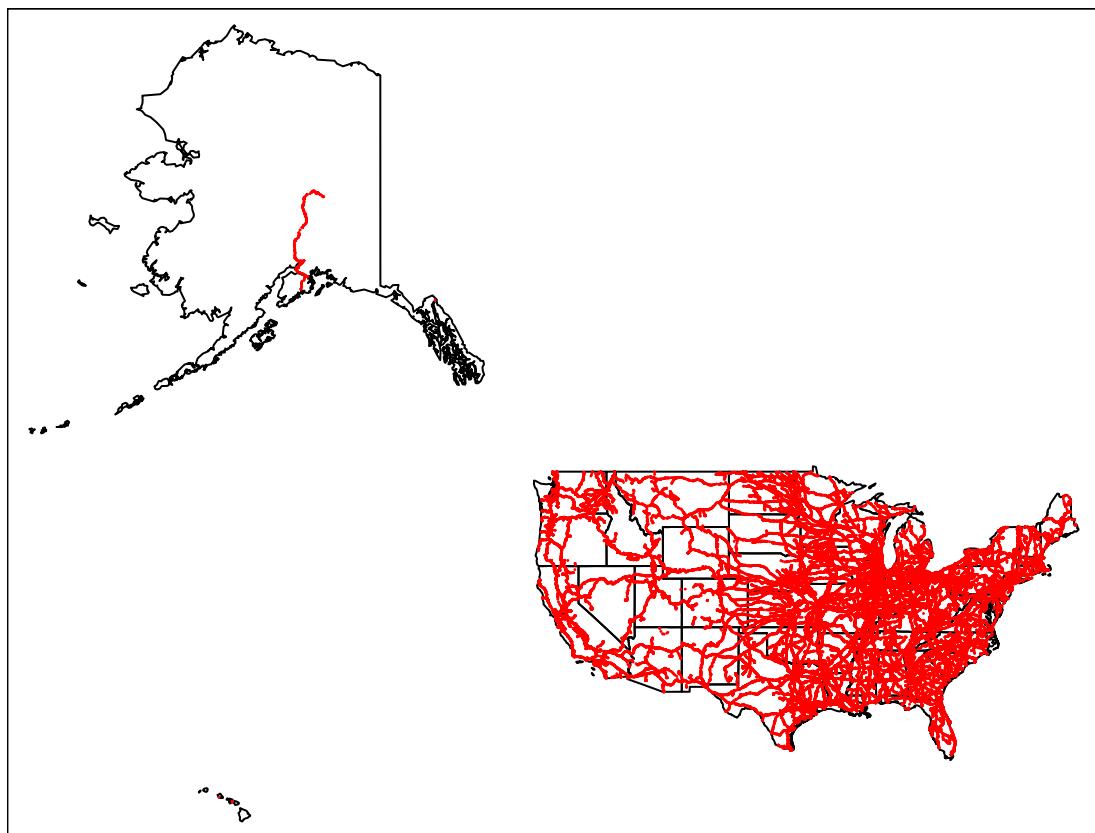
```
setattr("class","url")
request <- build_url(url)
USA_States_2017 <- st_read(request)

## Reading layer `OGRGeoJSON` from data source `https://services.arcgis.com/P3ePLMYs2RVChkJx/arcgis/res
## Simple feature collection with 51 features and 1 field
## geometry type:  MULTIPOLYGON
## dimension:      XY
## bbox:            xmin: -19839090 ymin: 2145730 xmax: -7454985 ymax: 11542620
## epsg (SRID):    3857
## proj4string:    +proj=merc +a=6378137 +b=6378137 +lat_ts=0.0 +lon_0=0.0 +x_0=0.0 +y_0=0 +k=1.0 +unit
library(tmap)

tmap_mode("plot")

## tmap mode set to plotting

tm_shape(USA_States_2017) +
  tm_borders("black") +
  tm_shape(exploited_railroads) +
  tm_symbols(col = "red", scale = 0.1, border.lwd = NA) +
  tm_legend(show = FALSE)
```



5 Statistics Netherlands (CBS) StatLine databank as open data

All tables in the [Statistics Netherlands \(CBS\) StatLine databank](#) are available as open data. Since 2014 this databank has an open data web API based on the OData protocol (<https://www.odata.org/>).

Using web services, data can be retrieved, filtered and combined. In this way, Statistics Netherlands aims to promote the widespread use of its statistical data.

In this chapter we will load Dutch statistical data directly into R,



5.1 the package cbsodataR

The package `cbsodataR` can be found on CRAN: <https://cran.r-project.org/package=cbsodataR>

The table of contents can be retrieved with the function `cbs_get_toc()`. As we browse through the table of contents we do see a total of 4422 entries. For most datasets all the metadata is in Dutch, but for some of them this information is also available in English.

When we split the TOC by language we do see that currently 701 datasets are available with metadata in English. These ‘English’ datasets are actually copies of their Dutch equivalents. (Maybe overtime all tables will be available with a description in English?)

This is a wealth of information, and all these datasets can be directly loaded into R. In the next paragraph we will do some exercises with one particular table.

```
library(cbsodataR)

toc <- cbs_get_toc()
nrow(toc)

## [1] 4422

toc_nl <- cbs_get_toc(Language = "nl")
nrow(toc_nl)

## [1] 3721

toc_en <- cbs_get_toc(Language = "en")
nrow(toc_en)

## [1] 701
```

5.2 Exercises with regional statistical data

5.2.1 Preparing the data

The table we have chosen for the exercises in this paragraph is called 'Regionale kerncijfers Nederland', i.e. regional statistical data about the Netherlands. The identifier of this table is **70072ned**. It is a huge table containing data at 5 regional levels - from the country as a whole down to the municipal level - from 1995 up to the current year.

The table can be retrieved with the function `cbs_get_data()`:

```
NL_Regional_Statistics2017 <- cbs_get_data('70072ned')
```

As we do not need the full table, we will limit the amount of data loaded.

Firstly, we are only interested in the year 2017. We want to analyse the total population by region for that particular year.

So, we come up with the following statement (we will explain the selection of columns below):

```
library(dplyr)
library(data.table)
library(stringr)

NL_Regional_Statistics2017 <- cbs_get_data('70072ned', Perioden = "2017JJ00") %>%
  select(RegioS, Perioden, TotaleBevolking_1, Code_291, Naam_292, Code_293, Naam_294)
```

The first two columns - `RegioS` and `Perioden` - are categorical columns, i.e. they contain codes. The labels for these columns can be added with the function `cbs_add_label_columns()`:

```
NL_Regional_Statistics2017 <- cbs_add_label_columns(NL_Regional_Statistics2017)
```

We will now rename the columns of our table for further use in the exercise:

```
NL_Regional_Statistics2017 <- rename(NL_Regional_Statistics2017,
  Code = RegioS,
  Name = RegioS_label,
  Year = Perioden,
  Year_label = Perioden_label,
  Total_population = TotaleBevolking_1,
  Region_code = Code_291,
  Region_name = Naam_292,
  Province_code = Code_293,
  Province_name = Naam_294)
```

Each column has a label. To avoid issues later on in this exercise, we will remove these labels now:

```
cols <- colnames(NL_Regional_Statistics2017)
for (c in cols) {
  attr(NL_Regional_Statistics2017[[c]], "label") <- NULL
}
```

Let's have a first look at the table:

```
head(NL_Regional_Statistics2017)
```

```
## # A tibble: 6 x 9
##   Code   Name   Year   Year_label Total_population Region_code Region_name
##   <fct> <fct> <fct> <fct>          <int> <fct>      <fct>
## 1 "NL0~ Nede~ 2017~ 2017           17081507 "        ~ "        ~
```

```

## 2 "LD0~ Noor~ 2017~ 2017          1722247 "
## 3 "LD0~ Oost~ 2017~ 2017          3603406 "
## 4 "LD0~ West~ 2017~ 2017          8125777 "
## 5 "LD0~ Zuid~ 2017~ 2017          3630077 "
## 6 "PV2~ Gron~ 2017~ 2017          583581 "LD01 ~ "Noord-Ned~
## # ... with 2 more variables: Province_code <fct>, Province_name <fct>

```

In the first row we can see the total number of inhabitants in the Netherlands on January 1st, 2017, which is: **17,081,507**.

And wait, we notice something else: some columns seem to have a fixed width. For example, the columns Code and Region_name:

```

paste(NL_Regional_Statistics2017$Code[1])
## [1] "NL01"
paste(NL_Regional_Statistics2017$Region_name[6])
## [1] "Noord-Nederland"
my_string <- paste(NL_Regional_Statistics2017$Region_name[6])
nchar(my_string)
## [1] 50

```

This is weird, isn't it? The string 'Noord-Nederland' - which is only 15 characters in length - is filled out with trailing spaces to have a length of 50... This can't be good, can it? We want our strings to behave like strings with their proper length, not carrying around a trail of whitespaces.

Let's fix this issue with the function trimws() for the columns concerned:

```

cols <- c('Code', 'Region_code', 'Region_name', 'Province_code', 'Province_name')
NL_Regional_Statistics2017[cols] <-
  lapply(NL_Regional_Statistics2017[cols], function(x){as.factor(trimws(x))})

head(NL_Regional_Statistics2017)

## # A tibble: 6 x 9
##   Code  Name  Year  Year_label Total_population Region_code Region_name
##   <fct> <fct> <fct> <fct>           <int> <fct>      <fct>
## 1 NL01 Nede~ 2017~ 2017          17081507   ""        ""
## 2 LD01 Noor~ 2017~ 2017          1722247    ""        ""
## 3 LD02 Oost~ 2017~ 2017          3603406    ""        ""
## 4 LD03 West~ 2017~ 2017          8125777    ""        ""
## 5 LD04 Zuid~ 2017~ 2017          3630077    ""        ""
## 6 PV20 Gron~ 2017~ 2017          583581 LD01     Noord-Nede~
## # ... with 2 more variables: Province_code <fct>, Province_name <fct>

```

This looks much better!

NUTS - Nomenclature des Unités Territoriales Statistiques

Eurostat uses a set of geospatial classifications, but the heart of these is NUTS: the **Nomenclature of Territorial Units for Statistics**.

This classification was set up by Eurostat at the beginning of the 1970s. It serves as a single, coherent system for dividing up the EU's territory in order to produce regional statistics for the European Union.

There are three levels of NUTS regions:

- NUTS 1: major socio-economic regions
- NUTS 2: basic regions for the application of regional policies
- NUTS 3: small regions for specific diagnoses



More information on the NUTS classification can be found here: <http://ec.europa.eu/eurostat/web/nuts/background>

Eurostat even created a nice video on the topic: <https://youtu.be/a4Y-hCQ-Klo>

5.2.2 Extracting NUTS 1 and NUTS 2 regions

Now that we have prepared our data, we can extract the data about NUTS 1 and NUTS 2 respectively. At NUTS 1 level the Netherlands is divided into 4 regions ('Landsdelen') and at NUTS 2 level into 12 provinces:

```
NL_Regions2017_data <- filter(NL_Regional_Statistics2017, Code %like% "LD")
NL_Provinces2017_data <- filter(NL_Regional_Statistics2017, Code %like% "PV")
```

Please note: these new data frames inherit factor levels from the original data frame from which they are created. Use the **RStudio Environment pane**, to confirm that - for example - the Name column of the NL_Provinces2017_data table is now a Factor with 770 levels, whereas there are only 12 provinces. So you have to drop unused levels before you can use the new subsets for further analysis:

```
# Make sure to drop unused levels from the factors in those new data.frames
NL_Regions2017_data <- droplevels(NL_Regions2017_data)
NL_Provinces2017_data <- droplevels(NL_Provinces2017_data)
```

```
NL_Provinces2017_data %>%
  select(Code, Name, Total_population, Region_code, Region_name) %>% as.data.frame()

##   Code           Name Total_population Region_code    Region_name
## 1 PV20    Groningen (PV)      583581     LD01 Noord-Nederland
## 2 PV21    Friesland (PV)      646874     LD01 Noord-Nederland
## 3 PV22     Drenthe (PV)      491792     LD01 Noord-Nederland
## 4 PV23 Overijssel (PV)      1147687     LD02 Oost-Nederland
## 5 PV24    Flevoland (PV)      407818     LD02 Oost-Nederland
## 6 PV25   Gelderland (PV)      2047901     LD02 Oost-Nederland
## 7 PV26     Utrecht (PV)      1284504     LD03 West-Nederland
## 8 PV27 Noord-Holland (PV)      2809483     LD03 West-Nederland
## 9 PV28  Zuid-Holland (PV)      3650222     LD03 West-Nederland
## 10 PV29     Zeeland (PV)      381568     LD03 West-Nederland
## 11 PV30 Noord-Brabant (PV)      2512531     LD04 Zuid-Nederland
## 12 PV31     Limburg (PV)      1117546     LD04 Zuid-Nederland
```

We have noticed that the province names in the column `Name` all have the suffix `(PV)`. Let's get rid of this suffix, while making sure that the column remains a factor:

```
NL_Provinces2017_data$Name <-
  as.factor(str_replace(NL_Provinces2017_data$Name, " \\\(PV\\\\)", ""))
```

Likewise, the region names in the column `Name` all have the suffix `(LD)`. We will also remove that:

```
NL_Regions2017_data %>% select(Code, Name, Total_population) %>% as.data.frame()
```

```
##   Code           Name Total_population
## 1 LD01 Noord-Nederland (LD)      1722247
## 2 LD02 Oost-Nederland (LD)      3603406
## 3 LD03 West-Nederland (LD)      8125777
## 4 LD04 Zuid-Nederland (LD)      3630077
```

```
NL_Regions2017_data$Name <-
  as.factor(str_replace(NL_Regions2017_data$Name, " \\\(LD\\\\)", ""))
```

5.2.3 barplot()

Let's investigate the spread of the Dutch population over the different regions by creating a barplot.


The most basic call to the `barplot()` function, is just to provide a column from a dataset and leave all the defaults:

```
barplot(NL_Provinces2017_data$Total_population)
# Try this yourself - result not printed in this manual
```

This returns a very basic plot, with no labels and no title, so not very useful.

Use `?barplot()` to see what arguments we can use. After some trial and error we might come to a statement like this:

```
barplot(NL_Provinces2017_data$Total_population / 1000000,
       names = NL_Provinces2017_data>Name,
       las = 2, cex.axis = .6, cex.names = .6,
       cex.main = .8, cex.lab = .8, space = 0,
       col = "lightblue",
       ylab = "Inhabitants (* 1.000.000)",
       main = "Number of Inhabitants by Province - The Netherlands - 2017")
```



Not bad at all, for a first attempt. Please note: the provinces are not presented in alphabetical order, but in their *natural* order, from Groningen (PV20) in the North to Limburg (PV31) in the South.

The provinces at the NUTS 2 level are grouped into regions ('Landsdelen') at the NUTS 1 level. We do know for each province to which region it belongs. Wouldn't it be nice to reflect this division in the plot?

So, In the next plot we will color the bars by region - North, East, West, South - using the `col` argument.

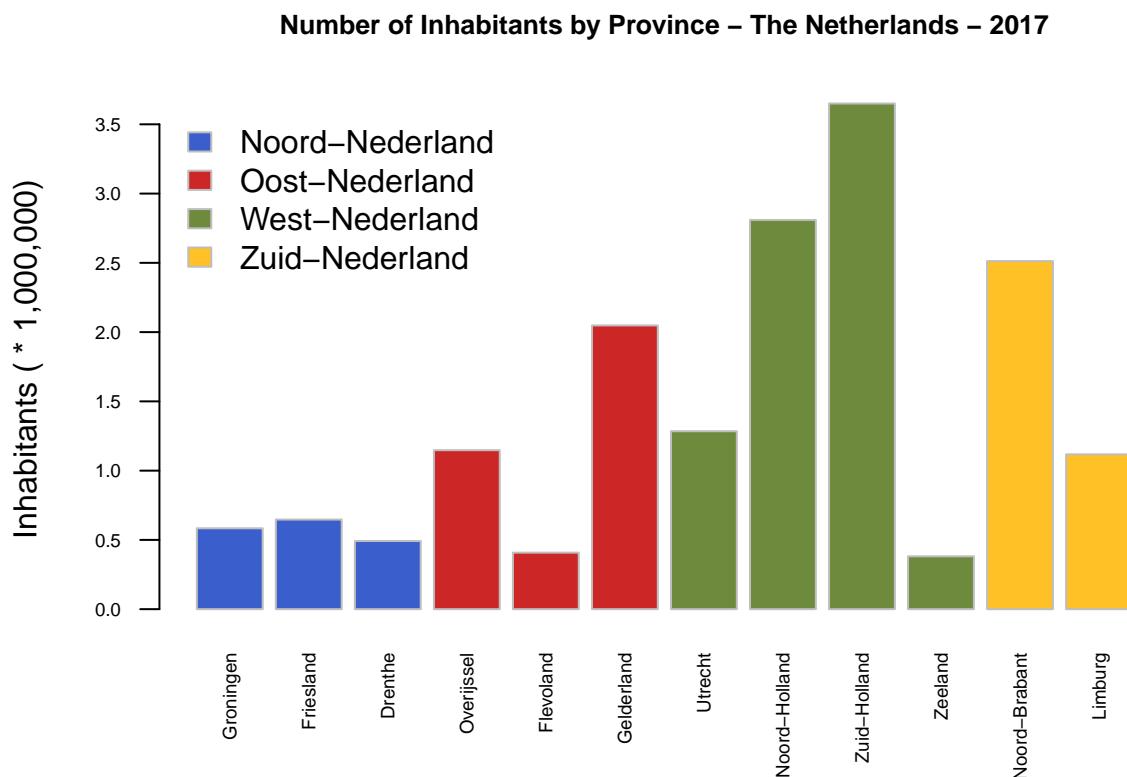
We will need 4 different colors for this plot. In the statement below we use the `palette()` function to manipulate the color

palette which is used when a col= has a numeric index. Using the c() function you create a vector with 4 distinct colors, which is passed as an argument into palette():

```
palette(c("royalblue3", "firebrick3", "darkolivegreen4", "goldenrod1"))
```

And now we can create the barplot like this - with the arguments for the legend provided as a list:

```
barplot(NL_Provinces2017_data$Total_population / 1000000,
        names = NL_Provinces2017_data>Name,
        las = 2, cex.axis = .6, cex.names = .6,
        cex.main = .8, border = "grey",
        col = NL_Provinces2017_data$Region_code,
        ylab = "Inhabitants (* 1,000,000)",
        main = "Number of Inhabitants by Province - The Netherlands - 2017",
        legend.text = unique(NL_Provinces2017_data$Region_name),
        args.legend = list(x = 'topleft',
                           bty = 'n',
                           fill = unique(NL_Provinces2017_data$Region_name),
                           border = 'grey'))
```



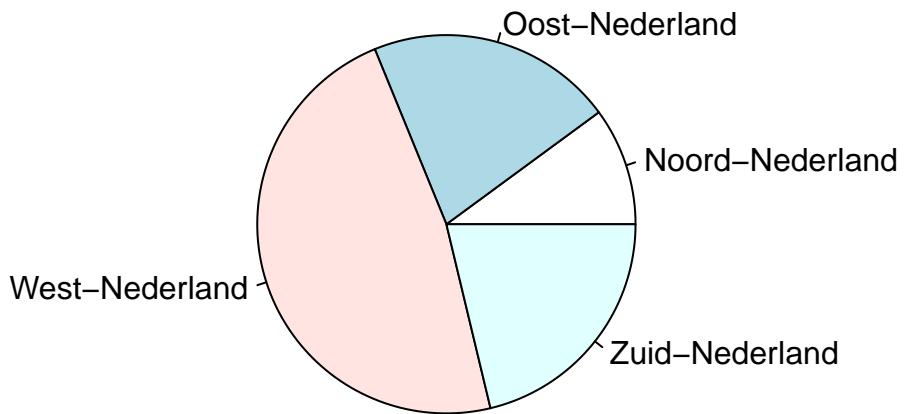
5.2.4 pie()

Based on the plots in the previous sections we might think that almost half the Dutch population lives in the westernmost region.

To confirm this hypothesis we will create a pie chart with the population by region.

This basic statement will create a simple pie chart:

```
pie(NL_Regions2017_data$Total_population, labels = NL_Regions2017_data>Name)
```

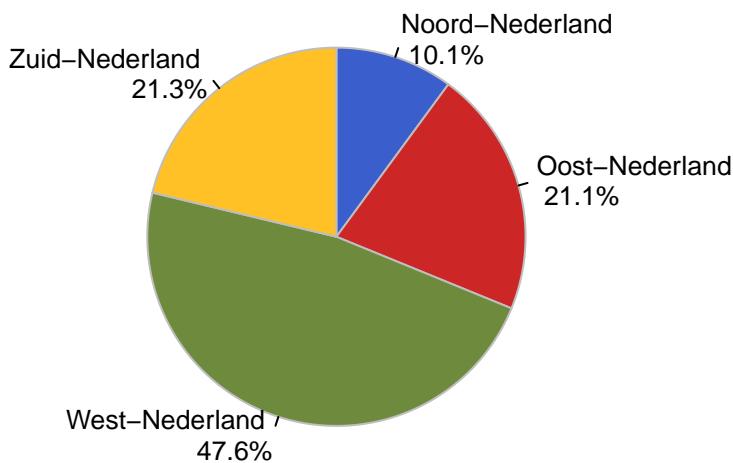


And yes, almost half of the Dutch population lives in West-Nederland.

The pie chart on the previous page already confirmed what we wanted to know. But of course - with a little extra effort - we can improve the quality of the graph by adding a title, calculating the actual percentages and using the same colors for the regions as we did in the barplot:

```
pct <- paste0(round(NL_Regions2017_data$Total_population /
                     sum(NL_Regions2017_data$Total_population) * 100, 1), "%")
lbls <- paste(NL_Regions2017_data>Name, "\n", pct)
palette(c("royalblue3", "firebrick3", "darkolivegreen4", "goldenrod1"))
pie(NL_Regions2017_data$Total_population, labels = lbls, clockwise = TRUE,
    cex = .8, col = NL_Regions2017_data>Name, border = "grey",
    main = "Percentage of Inhabitants by Region – The Netherlands – 2017")
```

Percentage of Inhabitants by Region – The Netherlands – 2017



5.2.5 Extracting municipal data

Now we will extract the municipalities from NL_Regional_Statistics2017.

```
NL_Municipalities2017_data <- filter(NL_Regional_Statistics2017, Code %like% "GM")
nrow(NL_Municipalities2017_data)
```

```
## [1] 713
```

Are there 713 municipalities in the Netherlands? No, not anymore! There used to be much more: in 1850 there were more or less 1200 municipalities. But today there are much less: the current (2018) amount is 380. The Netherlands has a long standing tradition of grouping smaller municipalities into larger ones. This is a slow process, but each year on the 1st of January the number of municipalities diminishes again.

We are looking at 2017, and in that year there were still 388 municipalities.

The table we have accessed contains data from 1995 up to today. And apparently in that year there were still 713 municipalities.

We can filter out the municipalities that do no longer exist by removing the ones without data. Here we use the column Total_population to check for that:

```
NL_Municipalities2017_data <- filter(NL_Municipalities2017_data, Total_population != "")
nrow(NL_Municipalities2017_data)
```

```
## [1] 388
```

And of course we should also drop unused factor levels:

```
NL_Municipalities2017_data <- droplevels(NL_Municipalities2017_data)
```

5.2.6 Merging sf and data.frame objects

Now we want to plot a map with the regional subdivision. To be able to do so we have to merge our data frame with a dataset containing geometry.

We will access a WFS server to download the municipal geometry - for the year 2017! - as we have done in paragraph 3.4.1:

```
# NL: Example with Dutch data
library(sf)
library(tmap)
library(httr)
library(data.table)
library(dplyr)

url <- list(hostname = "geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs",
            scheme = "https",
            query = list(service = "WFS",
                         version = "2.0.0",
                         request = "GetFeature",
                         typename =
                           "cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd",
                         outputFormat = "application/json")) %>%
  setattr("class", "url")
request <- build_url(url)

NL_Municipalities2017 <- st_read(request)

## Reading layer `OGRGeoJSON` from data source `https://geodata.nationaalgeoregister.nl/cbsgebiedsindelingen/wfs?service=WFS&version=2.0.0&request=GetFeature&typename=cbsgebiedsindelingen:cbs_gemeente_2017_gegeneraliseerd&outputFormat=application/json`
## Simple feature collection with 388 features and 5 fields
## geometry type: MULTIPOLYGON
## dimension: XY
## bbox: xmin: 13565.4 ymin: 306846.9 xmax: 277992.8 ymax: 619291
## epsg (SRID): 28992
## proj4string: +proj=sterea +lat_0=52.15616055555555 +lon_0=5.38763888888889 +k=0.9999079 +x_0=1550000 +y_0=406857,0.350733,-1.87035,4.0812 +units=m +no_defs
```

As we do not need all attribute columns, we will select only the relevant ones:

```
NL_Municipalities2017 <- select(NL_Municipalities2017, statcode, statnaam)
```

Please note: in the statement above we only select 2 columns, but the resulting object does also contain the geometry column. This is due to the *sticky* character of this geometry column. When you select a subset of columns, the geometry will always be included, unless you explicitly drop it. In other words: a subset of an sf object will in itself also be an sf object.

Now we will rename the columns:

```
NL_Municipalities2017 <- rename(NL_Municipalities2017, Code = statcode, Name = statnaam)
```

Both datasets do contain a column with municipal codes, so now you can merge them, like you would do with any two data frames:

```
NL_Municipalities2017 <-
  merge(NL_Municipalities2017, NL_Municipalities2017_data, by = "Code")

class(NL_Municipalities2017)

## [1] "sf"           "data.frame"
```

5.2.7 qtm()

And now we can plot our map with NUTS 1 regions, again using the same colors as before:

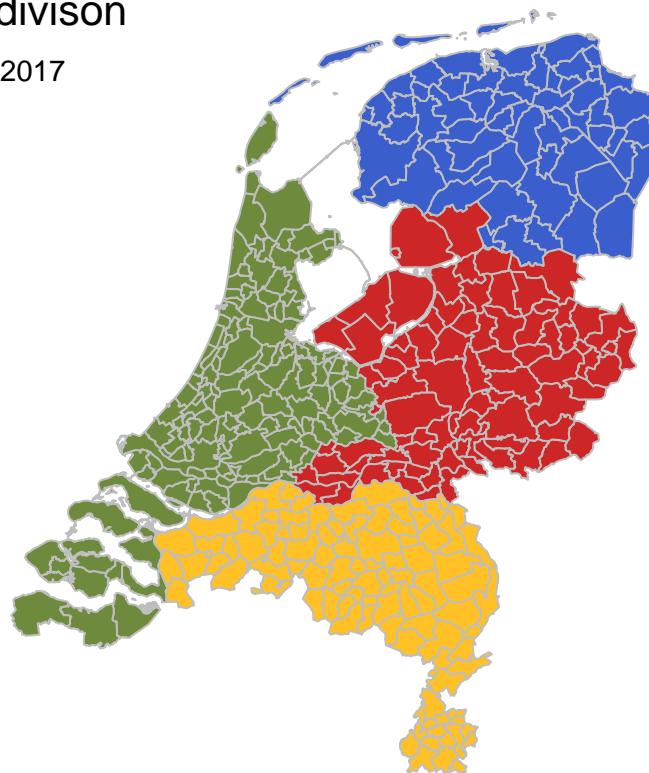
```
qtm(shp = NL_Municipalities2017,
    title = "Regional subdivision",
    fill = "Region_name",
    fill.title = "The Netherlands – 2017\nRegion",
    fill.palette =
      palette(c("royalblue3", "firebrick3", "darkolivegreen4", "goldenrod1")),
    borders = "grey",
    format = "NLD_wide")
```

Regional subdivision

The Netherlands – 2017

Region

- █ Noord–Nederland
- █ Oost–Nederland
- █ West–Nederland
- █ Zuid–Nederland



Factors in R

Conceptually, factors are variables in R which can only contain a pre-defined set of values, known as levels. Factors are often referred to as *categorical variables*, which can either be **ordered** (e.g. 'low', 'medium', 'high') or **unordered** (e.g. 'male', 'female').

One of the most important uses of factors is in statistical analysis and plotting. Storing categorical variables as factors insures that the statistical modeling functions will treat such data correctly. Factors are stored in R as a vector of integers, with a corresponding set of labels. These labels - one for each level - are used when the factor is displayed.

When data are being loaded into R, for example with the function `read.csv()`, all columns containing text (character data) will be automatically converted to factors. While this is very useful in many cases, there might also be situations where you will want to treat your data as continuous variables.

5.3 Factor: group municipalities into categories

In a previous exercise we have learned that there is an ongoing process of municipal regrouping in the Netherlands. Let's investigate opportunities for future regroupings by mapping municipalities by size.

We will start this exercise by dividing our municipalities in 3 groups: small, medium and large.

A small municipality has less than 50,000 inhabitants, a medium sized one has between 50,000 and 200,000, whereas a municipality with more than 200,000 inhabitants is considered to be large.

```
NL_Municipalities2017 <- mutate(NL_Municipalities2017,
  Category = case_when(Total_population < 50000 ~ "Small population",
    Total_population >= 50000 &
    Total_population < 200000 ~ "Medium population",
    Total_population >= 200000 ~ "Large population"))
```

```
class(NL_Municipalities2017$Category)
```

```
## [1] "character"
```

Now we will create an ordered factor with 3 levels:

```
population_levels <- c("Small population", "Medium population", "Large population")
NL_Municipalities2017 <- mutate(NL_Municipalities2017,
  Category = factor(Category, levels = population_levels, ordered = TRUE))
```

```
class(NL_Municipalities2017$Category)
```

```
## [1] "ordered" "factor"
```

And now we can plot our map.

Please note: in the map printed below there are not 3 categories, but 4. We have added a category 'very small' for municipalities with less than 20,000 inhabitants.

Question:

What steps do you need to take to add a new factor level after factors have been defined?

As you can see there is still some potential to cluster very small municipalities together. But of course, this will be more a political than a scientific issue.


```
qtm(shp = NL_Municipalities2017, title = "Municipalities by Population Size",
  fill = "Category", fill.title = "The Netherlands - 2017\nCategory",
  borders = "grey", format = "NLD_wide")
```

Municipalities by Population Size

The Netherlands – 2017

Category

- █ Very small population
- █ Small population
- █ Medium population
- █ Large population



6 Spatial Reference Systems

In this chapter we will have a quick look at an important topic when handling geospatial data, i.e. the Coordinate Reference System (or Spatial Reference System), which defines how spatial elements relate to the surface of the Earth and which is used to project our geodata on a flat screen. We have ignored this subject more or less in the previous exercises, mainly because of the fact that the whole matter of projecting the data is handled very well in the ‘sf’ package, or rather the ‘proj.4’ software behind it. We just didn’t have to think about it. Or wait, we did set a ‘crs’ explicitly - using an ‘epsg’ code - in the very first exercise this morning. Remember?

In general when you read in spatial data it comes with a proper

CRS and if all your data is in this same projection, you really don’t have to worry about all the geodesy behind it. But what if you want to combine data from different sources which appear to have different spatial reference systems?

Therefor we will shortly touch upon:

- Geographic coordinate systems vs. Projected coordinate systems
- the two main ways to describe a CRS in R: the ‘epsg’ code and the ‘proj4string’ definition
- how to reproject your data - temporarily or permanently

6.1 Coordinate systems: Geographic vs. Projected

Geographic coordinate systems identify locations on the Earth’s surface using longitude and latitude. Longitude is the distance East or West in angular distance, i.e. measured in degrees, from the Prime Meridian plane. Latitude is this distance North or South of the equatorial plane. And no, there is not just one ‘Lat\Lon coordinate system’. All geographic systems have to use the same reference, i.e. the Prime Meridian and the Equator, but the exact location on the Earth depends on several parameters like the ellipsoid and the datum used. So even if someone gives you data ‘in degrees’ you still have to find out which CRS to use. With a wrong CRS your data might end up centimeters or even several meters from the correct location. An important advantage of these geographic systems is their global reach - they can be used everywhere and are useful when dealing with international datasets. An important disadvantage is: distances cannot be easily measured in for example meters or kilometers in geographic CRSs.

Projected coordinate system have a much smaller reach, as these are based on Cartesian coordinates on an implicitly flat surface. They are only valid for a specific area - a country or even just a part of a country - where they have an origin, x and y axes, and a linear unit of measurement such as meters. We have seen several examples of data in such local, projected systems, e.g. EPSG:31287 for Austria, EPSG:3067 for Finland and EPSG:28992 for the Netherlands (but only onshore). These systems are ideal for dealing with data at a national or subnational level. In such a case you shuld really stick with the local CRS.

6.2 epsg codes and. proj4string definitions

As there are so many coordinate systems worldwide it is good to be clear about what system you use in your analysis. There are two main ways to describe a CRS in R: the epsg code and the proj4string definition. If you get your spatial data from a proper source, the `st_read()` function will return both, as we have seen in the exercises before.

The epsg code gets its name from the - no longer existing - *European Petroleum Survey Group*. The set of EPSG definitions is currently maintained by the International Association of Oil & Gas Producers (IOGP). Apparently location is very important in the fossil fuel industry.

An epsg code refers to only one, well-defined coordinate reference system. This coding system is widely adopted by the geospatial community, so if you get XY data with an existing epsg code you can easily plot your data on the map, in R or in any other GIS system for that matter.

A proj4string definition contains different parameters such as the projection type, the datum and the ellipsoid. It allows you to specify your own projection, or to modify an existing one. But in general it would be advisable to refer to the data provider if there is something wrong with the projection of your data.

6.3 Reprojecting data

References

- Lovelace, Robin, Jakub Nowosad & Jannes Muenchow (Forthcoming). *Geocomputation with R*. CRC Press. Online version: <https://geocompr.robinlovelace.net/>
- Pebesma, Edzer (January 3, 2017). *Simple Features Now on CRAN*. Blog. Retrieved from: <https://www.r-consortium.org/blog/2017/01/03/simple-features-now-on-cran>

Appendix A: List of abbreviations used

Abbreviation	Meaning
API	application programming interface
CRAN	Comprehensive R Archive Network
GIS	Geographic Information System
GML	Geography Markup Language
IOGP	International Association of Oil & Gas Producers
JSON	JavaScript Object Notation
OGC	Open Geospatial Consortium
REST	Representational State Transfer
SaaS	software as a service
WFS	Web Feature Service
XML	eXtensible Markup Language

Appendix B: Additional exercises

OGC Web Feature Service (WFS)

For background information on the steps taken in the exercise below, please refer to chapter 3

request=GetFeature - another example from the Netherlands

In this additional exercise we will access a WFS service offering information about *bevolkingskernen* (human settlements) in the Netherlands in 2011. More information on this dataset can be found here: [Dataset: CBS Bevolkingskernen](#).

The URL with GetCapabilities request is:

<https://geodata.nationaalgeoregister.nl/bevolkingskernen2011/wfs?request=GetCapabilities>

In the <FeatureTypeList> section of this XML response we can see that this service only offers one <FeatureType>, i.e. one dataset. This dataset has the <Name> **bevolkingskernen2011:cbsbevolkingskernen2011**. That's the value we are giving to the typename parameter in our GetFeature request.

The URL with DescribeFeatureType request is:

<https://geodata.nationaalgeoregister.nl/bevolkingskernen2011/wfs?service=WFS&version=2.0.0&request=DescribeFeatureType&typename=bevolkingskernen2011:cbsbevolkingskernen2011>

This dataset does contain quite some attribute columns!

This is the script to build the full GetFeature request:

```
# NL: Human Settlement Analysis
library(sf)
library(tmap)
library(httr)
library(data.table)

url <- list(hostname = "geodata.nationaalgeoregister.nl/bevolkingskernen2011/wfs",
            scheme = "https",
            query = list(service = "WFS",
                         version = "2.0.0",
                         request = "GetFeature",
                         typename = "bevolkingskernen2011:cbsbevolkingskernen2011",
                         outputFormat = "application/json")) %>%
  setattr("class", "url")
request <- build_url(url)
```

The variable `request` will now contain this [link](#). (When you click this link the server response will be shown in your browser in GeoJSON format.)

Now we want to feed this response directly into R, like this:

```
NL_Human_Settlements2011 <- st_read(request)

## Reading layer `OGRGeoJSON` from data source `https://geodata.nationaalgeoregister.nl/bevolkingskernen2011/wfs?request=GetFeature&typename=bevolkingskernen2011:cbsbevolkingskernen2011&outputFormat=application/json`
## Simple feature collection with 2168 features and 112 fields
## geometry type:  POLYGON
## dimension:      XY
## bbox:           xmin: 15153.79 ymin: 308078.3 xmax: 277312.9 ymax: 610955.6
## epsg (SRID):   28992
## proj4string:   +proj=sterea +lat_0=52.15616055555555 +lon_0=5.387638888888889 +k=0.9999079 +x_0=1550000 +y_0=660000
```

```
plot(st_geometry(NL_Human_Settlements2011), col = "orange", border = "red")
```

